

# Transcript: Meeting Max XAI

May 9, 2026 | Rick de Rijk and Max | Duration: 24:00 | Language: English

---

## TASK

---

### [00:08] Rick de Rijk

This is essentially a counterfactual task. It means: how can you modify something so that a five and a six look alike. That is roughly what this will demonstrate.

### [00:19] Max

So the AI will think it is a six, but it is actually nothing.

### [00:24] Rick de Rijk

Right, for a human it is interpretable as nothing, but you will see that shortly. You can take a moment to read through the protocol first. There are also two different metrics: validity and plausibility. There is also an explanation of what those mean exactly. When you go through the dashboard, just say out loud what you see, what you think of it, whether it is clear or not.

### [01:19] Rick de Rijk

You have the validity and plausibility metrics here. They are explained with a few examples of how you can measure them. When you are done I have five closing questions.

### [01:33] Max

So a visual inspection.

### [01:39] Rick de Rijk

Yes, it is a guided approach. The idea is that users and data science students go through the steps one by one to understand what the model is doing, and what the difference is between human judgment and model output.

### [02:46] Rick de Rijk

Now you can see the dashboard. At the top there is an information bar explaining what the task is about. You can just work through it step by step.

### [02:49] Max

Okay, let me have a look. This is a nice layout.

### [04:22] Max

I think this one is a 'certain'.

### [04:28] Rick de Rijk

Then you should also be able to describe why you think that. English or Dutch, does not matter.

### [05:08] Rick de Rijk

At the top there are information blocks, so you can always scroll back up to re-read what something was.

**[05:21] Max**

The model has classified this as a 1. This one then has a 0. So the one is not a 5.

**[06:13] Rick de Rijk**

You have five methods. If you scroll down you will see all five of them.

**[06:18] Max**

And these are the images the model is looking at?

**[06:20] Rick de Rijk**

Yes. You now need to decide which one you think is the best out of the methods you see, and how confident you are in that choice. So what is best from your perspective, and that will then be compared to what the model found to be best.

**[06:41] Max**

Let me see, that is the Alibi Proto CF. I will go with that one. I can see a seven here.

**[07:01] Rick de Rijk**

Here you can explain why you think that.

**[08:42] Rick de Rijk**

You can also just say that something is unclear. It is also meant to check whether you find it clear.

**[09:02] Rick de Rijk**

It shows how confident the model is that it is a 7. Green means the score is above average. Red means below average. The idea is to make it visible how good the metric score is.

**[09:27] Max**

Right, for some metrics it is red.

**[09:34] Max**

On average I would say: borderline. This one is green, this one is red.

**[11:25] Rick de Rijk**

You have now seen which method is best according to the dashboard.

**[11:28] Max**

Yes, according to the model itself.

**[11:29] Rick de Rijk**

Yes. And here you can see that the counterfactual was produced by MinEdit.

**[11:38] Max**

I thought the Alibi Proto. That was the one I had chosen.

**[11:41] Rick de Rijk**

Yes, and then you can also see that that was your choice.

**[11:45] Max**

The lower the M1 score, the better the image. Let me see: 'Look at the foreground in bold. Does it make sense to you?'

**[12:02] Rick de Rijk**

Then you can say: yes, no, or no because I chose this method, et cetera.

**[13:10] Max**

'How do you think this explains the difference between human judgement and the metrics?' That sounds like a good answer.

**[13:43] Rick de Rijk**

This is a summary of what you clicked earlier, as a reminder. And here is an additional explanation.

**[14:06] Max**

Something shorter would be nicer.

**[14:09] Max**

Or you would want the metrics to be shown here with an explanation next to them.

**[14:17] Rick de Rijk**

Yes, per block essentially.

**[14:19] Max**

Or that you can hover over it and then get an explanation.

**[14:24] Rick de Rijk**

So more of a user interface thing really.

**[14:28] Max**

Yes. If I want to know something now, I have to scroll up. And then by the time I scroll back down I have already forgotten it.

**[15:05] Rick de Rijk**

Are the metrics not an indicator for you? Do you look purely visually?

**[15:11] Max**

Yes, purely visual. The metrics do not really interest me.

**[15:17] Rick de Rijk**

You can also write that down: that you find the metrics completely uninformative, or that you do not understand what they add.

**[15:27] Max**

I just see a number, that is what I use.

**[15:31] Rick de Rijk**

So you reason from your own visual judgment.

**[15:33] Max**

Yes, from my own way of working.

**[16:34] Rick de Rijk**

This is the last section. When you are done here, do not click OK yet. I will ask the final question first and then you can use it. This ending is a summary of all your results.

**[17:13] Rick de Rijk**

This last section is purely a summary. You can come back to earlier choices if you want. Before you click Final: wait a moment, otherwise you will lose your input.

**[19:02] Rick de Rijk**

The model can also time out, meaning it cannot always turn a 3 into a 5. The idea is: the model adjusts certain pixels. Then the model thinks: a 3 has now become a 5. But people still think: I see nothing, this cannot be right at all. That is the concept.

---

## **POST-INTERVIEW**

---

### **Q1 – Decision Process**

**[19:40] Rick de Rijk**

Can you explain step by step why you found that method to be the best? That was the Alibi Proto.

**[20:00] Max**

You could very clearly see a digit in that method, and not in the others.

### **Q2 – Alignment with Metrics**

**[20:13] Rick de Rijk**

The metrics you saw, did they meet your expectations?

**[20:28] Max**

A model can produce those metrics, but if I still see nothing in the image, I do not assign any value to those metrics.

### **Q3 – Interface Understanding**

**[20:40] Rick de Rijk**

At what point did you understand the difference between your judgment and the metrics? What helped or did not help?

**[21:33] Max**

Those metrics are based on what the model sees. If the model is wrong at the front end but too confident that it is a digit, then those metrics still have value for the model.

**[21:45] Rick de Rijk**

So essentially the metrics did not help you at all?

**[21:49] Max**

No.

#### **Q4 – Dashboard Improvement**

**[21:49] Rick de Rijk**

If you could change one thing in the dashboard, in terms of layout, order of steps, or how metrics are shown, what would it be?

**[22:06] Max**

Just make sure you have to scroll less.

**[22:14] Max**

Yes, that you can hover over a metric and then get a short explanation.

**[22:23] Max**

And also a bit less text. There is quite a lot of it right now.

**[22:50] Rick de Rijk**

Was the order of the steps logical? From step 1 to the end?

**[23:16] Max**

Yes, I think the order is fine.

**[23:19] Rick de Rijk**

Was there a moment where you were confused about what to do? That was at step 4 of the Action Matrix right?

**[23:31] Max**

Yes.

**[23:48] Max**

That was a bit unclear. But if you want to move through quickly, you also do not read everything.