

XAI Report

Prediction & Predictive coming together

Rick de Rijk, Stefan Vonk, Tycho van Rooij, Thom Verzantvoort & Gilbert Laanen

Table of contents

1	1 Introduction	3
1.1	1.1 Assignment Description	3
1.2	1.2 Design Debrief	4
2	2 Empathize Phase	4
2.1	2.1 Methods Used	4
2.1.1	2.1.1 Affinity Mapping Workshop	5
2.1.2	2.1.2 User Identification and Target Group Definition	6
2.1.3	2.1.3 Literature Scan	7
2.1.4	2.1.4 Technical Inspection of the Dataset	7
2.2	2.2 Results and Insights	7
2.2.1	2.2.1 Insight 1: Users struggle to distinguish validity from plausibility	7
2.2.2	2.2.2 Insight 2: Metrics can influence user judgment	7
2.2.3	2.2.3 Insight 3: Comparing multiple explanations is cognitively demanding	9
2.2.4	2.2.4 Insight 4: The gap between intuition and metrics is not explicit	9
2.2.5	2.2.5 Insight 5: Users benefit from structured guidance	9
2.3	2.3 Transition to Define Phase	10
3	3 Define Phase	10
3.1	3.1 Problem Statement	10
3.2	3.2 Target User	10
3.3	3.3 Focus for Ideation	11
4	4 Ideation Phase	11
4.1	4.1 Divergence — Generating Design Concepts	11
4.1.1	4.1.1 Free Exploration Dashboard	11
4.1.2	4.1.2 Metric-First Evaluation Interface	12
4.1.3	4.1.3 Explanation Comparison View	12
4.1.4	4.1.4 Ranking-Based Interface	14
4.1.5	4.1.5 Guided Task Mode	14
4.2	4.2 COCD Box — Structuring Divergence and Convergence	15
4.3	4.3 Convergence — Selection of Design Direction	16
4.3.1	4.3.1 Alignment with User Needs	16
4.3.2	4.3.2 Mitigating Metric Bias	16
4.3.3	4.3.3 Making the Core Problem Observable	16
4.4	4.4 Design Direction	16

5	Prototype Phase	17
5.1	Prototype Description	17
5.1.1	Code Structure in <code>src/</code>	17
5.1.2	Result Storage in <code>Game Results/</code>	17
5.2	Design Decisions	18
5.2.1	Linear, locked step progression	18
5.2.2	Showing one random CF method in steps 1 and 2 before the full lineup	18
5.2.3	Combining free-text and structured input	18
5.2.4	In-session guidance and explainability support	18
5.2.5	Hiding metrics until step 4	18
5.2.6	The Game as the primary data collection mechanism	19
5.2.7	Collecting all inputs per session, not per step	19
5.2.8	Dataset choice	19
5.2.9	Archive	19
6	Test Phase (4 pages)	20
6.1	Research Question	20
6.2	Method	20
6.3	Results	20
6.4	Insights	20
7	Conclusions & Recommendations (4 pages)	20
7.1	Conclusions from User Testing	20
7.2	Improvement Recommendations	20
7.3	Final Design Concept	20
7.4	User Interaction in Context	20
8	Design Archive Description (0.5 page)	20
8.1	Archive Overview	20
9	References	21
10	Appendix Final Version Prototype	21
11	Appendix Archive Prototype	29

List of Figures

1	Affinity mapping session showing early problem exploration and user insights	5
2	User groups identified during the workshop	6
3	Clustering of insights related to evaluation challenges	8
4	Notes highlighting confusion and bias in interpretation	8
5	Insights related to comparison and decision-making challenges	9
6	Brainwriting session showing a wide range of alternative dashboard concepts	12
7	Generated ideas focusing on interaction patterns and evaluation challenges	13
8	COCD box showing divergence and convergence of design ideas	15
9	HomeScreen 1	21
10	Homescreen 2	21
11	Visual Inspection FV	22
12	User Prediction FV	22
13	The Game	22
14	Actual metrics 1	23

15	Actual metric 2	23
16	Explanation & Feedback FV 1	24
17	Explanation & Feedback 2	24
18	Compare methods FV 1	25
19	Compare methods 2 FV	25
20	Compare methods FV 3	25
21	CF & Thoughts FV	26
22	CF Thoughts 2	26
23	Mini-Game Homescreeen	27
24	Minigame Information	27
25	Round Example 1	27
26	Round Example 2	28
27	Mini-Game Choice	28
28	Leaderbord	28
29	Homescreeen	29
30	Visual Inspection	29
31	User Prediction	30
32	The Game	30
33	Actual Metrics	31
34	Explanation & Feedback	31
35	Compare methods 1.1	31
36	Compare methods 1.2	32
37	Archive CF & FT	32

1 1 Introduction

1.1 1.1 Assignment Description

This project is conducted within the course *Interactive and Explainable AI Design* and focuses on Assignment B, which involves designing and evaluating an interactive dashboard for exploring counterfactual explanations. The assignment is situated in the domain of explainable artificial intelligence (XAI), where the goal is not only to provide explanations for model predictions, but also to understand how users interpret and interact with these explanations in practice (Guidotti et al. 2018).

Counterfactual explanations are a widely used post-hoc explanation technique, first formalized by (Wachter, Mittelstadt, and Russell 2017). A counterfactual describes the minimal change to an input instance required to obtain a different model prediction, effectively answering the question of what needs to change for a different outcome to occur. In the context of image classification, this corresponds to modifying an image such that the model predicts a specified target class instead of the original prediction.

A key challenge in this domain is the evaluation of counterfactual explanations. Multiple evaluation criteria exist, of which validity and plausibility are among the most prominent. Validity refers to whether the counterfactual successfully changes the model’s prediction to the desired target class, while plausibility captures whether the resulting explanation appears realistic or meaningful to a human observer.

Importantly, these evaluation dimensions do not always align. A counterfactual may be valid according to the model, yet appear implausible to a human, or vice versa. This reflects a broader issue in explainable AI, where objective evaluation metrics and human perception can diverge (Doshi-Velez and Kim 2017). Prior research suggests that different explanation methods and evaluation criteria may lead to conflicting interpretations, creating ambiguity for users attempting to assess explanation quality (Guidotti et al. 2018).

The objective of this assignment is therefore twofold. First, to design a dashboard that enables users to explore and compare counterfactual explanations in a structured and intuitive way. Second, to evalu-

ate how users interpret these explanations and whether their judgments align with objective evaluation metrics.

1.2 1.2 Design Debrief

The goal of this project is to design an interactive dashboard that reveals and helps analyse the gap between human intuition and model-based evaluation of counterfactual explanations. The primary target users are data science students, with machine learning researchers as a secondary audience. These users are assumed to have a basic understanding of model evaluation, but may still experience difficulty in interpreting counterfactual explanations consistently.

The core problem addressed in this design is that users tend to rely on visual intuition when evaluating explanations, while objective metrics such as validity and plausibility provide a different, and sometimes conflicting, perspective. Furthermore, prior work indicates that the way explanation information is presented can influence user judgment, for example by anchoring decisions on numerical metrics rather than independent reasoning (Kaur et al. 2020).

To address this issue, the project introduces a *Guided Task Mode* dashboard. Instead of supporting unrestricted exploration, the interface guides users through a structured sequence of evaluation steps. Users first assess counterfactual explanations based solely on visual inspection, then estimate evaluation metrics, and only afterwards are the actual metrics revealed. This controlled progression is intended to support independent judgment before exposing users to model-based evaluation, enabling a clearer comparison between intuition and objective metrics.

The dashboard functions both as an analysis tool and as a data collection instrument. It captures user judgments, metric estimates, method selections, and qualitative reflections within a single interaction flow. This makes it possible to analyse how users interpret counterfactual explanations, how confident they are in their judgments, and where discrepancies between human perception and model-based evaluation may occur.

This design follows a user-centered perspective on explainable AI, in which explanations are not only evaluated based on technical correctness, but also on how effectively they support human understanding and decision-making (Amershi et al. 2019).

Based on this motivation, the central research question guiding this project is:

i Research Question

To what extent do human judgments of counterfactual explanation quality align with objective evaluation metrics, and how can interface design reveal or reduce discrepancies between the two?

2 2 Empathize Phase

2.1 2.1 Methods Used

The goal of the empathize phase was to develop a well-grounded understanding of both the problem space and the intended user group. In line with the design thinking framework, this phase focuses on uncovering user needs, identifying pain points, and understanding how current approaches to counterfactual explanations may limit effective interpretation.

To achieve this, three complementary methods were applied: (1) a structured affinity mapping workshop, (2) a literature scan of explainable AI (XAI) challenges and tools, and (3) a technical inspection of the provided counterfactual dataset. These methods were combined to ensure both user-centered and technology-oriented insights.

2.1.1 Affinity Mapping Workshop

A collaborative workshop was conducted in which all team members independently generated observations, questions, and assumptions related to counterfactual explanations. These were written on post-it notes and subsequently grouped using affinity mapping. This approach allowed the team to externalize assumptions and identify patterns across perspectives.




Figure 1: Affinity mapping session showing early problem exploration and user insights

The workshop focused on three key areas:

1. understanding how users interpret counterfactual explanations
2. identifying sources of confusion or misinterpretation
3. exploring potential use cases and interaction patterns

The clustering process suggested several recurring themes. Many notes pointed to the difficulty of determining which counterfactual explanation is “better,” especially when multiple methods produce different results. Other clusters indicated uncertainty about how to interpret evaluation metrics and when these

Importantly, this method surfaced both preliminary insights and open questions, such as how users form judgments without guidance and whether they rely more on visual intuition or numerical metrics.

During the workshop, potential user groups were identified and refined. The primary user group was defined as data science students, while machine learning researchers were considered a secondary group.




Figure 2: User groups identified during the workshop

Data science students are expected to have a basic understanding of machine learning concepts but limited experience with evaluating explanation quality. They may therefore rely more strongly on visual intuition and may not fully understand the implications of different evaluation metrics.

Machine learning researchers, in contrast, are more familiar with model evaluation but may still face challenges when comparing multiple explanation methods or interpreting trade-offs between different metrics.

2.1.3 Literature Scan

A targeted literature scan was conducted to contextualize the observed challenges within existing research on explainable AI. Prior work indicates that explanation quality is inherently multi-dimensional and that no single metric fully captures what makes an explanation “good” (Guidotti et al. 2018).

Furthermore, research suggests that there can be a gap between objective evaluation metrics and human interpretability. Doshi-Velez and Kim argue that interpretability should be evaluated with respect to human understanding, rather than purely technical criteria (Doshi-Velez and Kim 2017).

In addition, studies on human interaction with explanation tools indicate that users can be influenced by how information is presented. For example, numerical metrics may lead to anchoring effects, where users base their judgment on the metric rather than forming an independent assessment (Kaur et al. 2020).

These findings provided a theoretical basis for interpreting the workshop results.

2.1.4 Technical Inspection of the Dataset

A technical inspection of the dataset was conducted by manually exploring a subset of cases in which multiple counterfactual explanations were generated for the same input instance. For each case, the outputs of different methods were compared alongside their associated evaluation metrics.

This inspection suggested that different methods can produce substantially different counterfactuals for the same input. In several observed cases, explanations that appeared visually plausible did not achieve the intended target prediction, while other explanations that were classified as valid by the model appeared less intuitive from a human perspective.

These observations indicate that trade-offs between different evaluation dimensions are present, and that explanation quality cannot be fully captured by a single criterion.

2.2 Results and Insights

The combination of methods resulted in a set of key insights that informed the subsequent design phases.

2.2.1 Insight 1: Users struggle to distinguish validity from plausibility

A central observation from the affinity mapping session is that users have difficulty distinguishing between explanations that are technically correct and those that appear visually convincing.

Users tend to interpret explanations based on visual similarity or perceived realism, rather than on whether the explanation actually changes the model prediction. As a result, explanations that are valid may be rejected, while implausible explanations may be accepted.

2.2.2 Insight 2: Metrics can influence user judgment

Another insight is that the presence of evaluation metrics can influence how users interpret explanations.





Figure 3: Clustering of insights related to evaluation challenges



Participants indicated that numerical values shown early in the process may guide decision-making. Instead of forming an independent judgment, users may rely on the metric as a heuristic, which can reduce critical reflection.

2.2.3 Insight 3: Comparing multiple explanations is cognitively demanding

Users expressed difficulty in comparing multiple counterfactual explanations simultaneously.




Figure 5: Insights related to comparison and decision-making challenges

The presence of multiple methods introduces several challenges:

- uncertainty about which criteria to use for comparison
- difficulty in tracking differences between explanations
- increased cognitive load when many options are presented

2.2.4 Insight 4: The gap between intuition and metrics is not explicit

An important insight is that the discrepancy between human intuition and model-based evaluation is not explicitly represented in existing tools.

Users may form an initial judgment based on visual inspection, but this judgment is typically not compared to objective metrics in a structured way. As a result, potential misalignment remains implicit.

2.2.5 Insight 5: Users benefit from structured guidance

A final insight is that users may benefit from structured guidance when interpreting explanations, rather than fully open exploration.

Without guidance, users may:

- focus on less relevant aspects of the explanation
- misinterpret evaluation metrics
- draw inconsistent conclusions

2.3 2.3 Transition to Define Phase

The empathize phase highlights a consistent pattern: users experience difficulty interpreting counterfactual explanations, particularly due to the misalignment between visual intuition and objective evaluation metrics. This challenge is reinforced by the way information is presented and by the complexity of comparing multiple explanations.

Based on these findings, the design focus is narrowed to the following challenge:

i How to design an interface that reveals and supports the interpretation of the gap between human intuition and model-based evaluation.

This forms the basis for the define phase, where the problem will be formalized into a clear problem statement.

3 3 Define Phase

3.1 3.1 Problem Statement

The empathize phase revealed a consistent challenge: users experience difficulty in interpreting counterfactual explanations due to the misalignment between visual intuition and objective evaluation metrics. While users tend to rely on visual similarity and perceived plausibility, model-based metrics such as validity may indicate a different assessment of explanation quality.

This leads to the following problem statement:

i *A data science student needs a way to systematically interpret and compare counterfactual explanations, because current representations do not make the relationship between human intuition and evaluation metrics explicit, resulting in confusion and inconsistent judgments.*

This problem is directly grounded in the insights from the empathize phase, where users expressed uncertainty about how to interpret metrics, difficulty in comparing multiple explanations, and a tendency to rely on intuition without validation.

3.2 3.2 Target User

The primary target user for this project is a data science student. These users typically have a foundational understanding of machine learning models and evaluation concepts, but limited experience with interpreting explanation quality, especially in situations where multiple evaluation criteria are involved.

Data science students are likely to:

- rely on visual intuition when evaluating explanations
- have difficulty interpreting abstract evaluation metrics
- struggle to compare multiple counterfactual methods

A secondary user group consists of machine learning researchers. While these users are more experienced with model evaluation, they may still face challenges when comparing explanation methods or analysing trade-offs between different evaluation criteria.

The design primarily focuses on supporting data science students, while remaining relevant for more advanced users.

3.3 Focus for Ideation

Based on the insights from the empathize phase, the design focus is narrowed to one central challenge: making the gap between human intuition and model-based evaluation visible and interpretable.

This focus is derived from three key observations:

1. Users struggle to distinguish between validity and plausibility
2. Evaluation metrics can influence or bias user judgment
3. Comparing multiple explanations introduces cognitive complexity

These observations indicate that the core issue is not only the availability of explanations, but how they are interpreted and compared.

Rather than designing a dashboard that prioritizes exploration or metric visualization, the design will focus on structuring the interaction in such a way that users first form an independent judgment, and only then are confronted with model-based evaluation.

This leads to the following design direction:

- guide users through a step-by-step evaluation process
- separate intuition from metric-based evaluation
- explicitly compare user judgments with model outputs

This direction forms the basis for the ideation phase, where multiple design concepts will be explored to determine how such a guided interaction can be implemented effectively.

4 Ideation Phase

4.1 Divergence — Generating Design Concepts

Following the define phase, the ideation phase focused on exploring a broad range of possible solutions to address the central challenge: *making the gap between human intuition and model-based evaluation visible and interpretable*.

To ensure sufficient divergence, a structured brainwriting session was conducted. Each team member independently generated ideas before sharing them with the group. This approach reduced the risk of early convergence on familiar solutions and encouraged the exploration of alternative interaction patterns.

Figure 6 illustrates the divergence phase. The variety of post-it notes reflects multiple directions, including dashboard-based exploration, metric-driven evaluation, comparison views, and more structured interaction flows.

Ideas were intentionally not evaluated during the initial generation phase. This allowed both conventional and more experimental concepts to emerge before introducing selection criteria.

The generated ideas can be grouped into several concept categories.

4.1.1 Free Exploration Dashboard

One concept focused on a traditional dashboard approach, where users can freely explore counterfactual explanations using filters, dropdowns, and coordinated visualizations.

Strengths:

- Supports flexible exploration and user control
- Aligns with familiar dashboard interaction patterns
- Allows users to define their own analysis process

Limitations:




Figure 6: Brainwriting session showing a wide range of alternative dashboard concepts

- Requires users to determine their own evaluation strategy
- Does not reduce cognitive load when comparing multiple explanations
- Offers limited support for interpreting evaluation metrics

Relation to the core problem: This concept does not explicitly address the gap between intuition and model-based evaluation. It assumes that users are able to interpret explanations independently, which contradicts the findings from the empathize phase.

4.1.2 Metric-First Evaluation Interface

Another concept prioritized evaluation metrics as the primary entry point. Users would first see values such as validity and plausibility, and then inspect the corresponding counterfactual explanations.

Strengths:

- Supports systematic and metric-driven evaluation
- Makes quantitative differences between methods explicit
- Efficient for expert users

Limitations:

- Risks influencing user judgment through anchoring effects
- Reduces the role of independent interpretation
- May discourage critical reflection

Relation to the core problem: This concept potentially reinforces the misalignment between intuition and metrics by prioritizing numerical values over user interpretation. It therefore does not support the goal of making this gap visible.

4.1.3 Explanation Comparison View

A third concept focused on side-by-side comparison of multiple counterfactual explanations.




Figure 7: Generated ideas focusing on interaction patterns and evaluation challenges

Strengths:

- Enables direct comparison between different methods
- Makes visual differences between explanations observable
- Supports exploratory analysis

Limitations:

- Introduces cognitive load when multiple explanations are presented
- Does not provide guidance on how to evaluate differences
- Leaves interpretation strategy undefined

Relation to the core problem:

While this concept supports comparison, it does not explicitly structure the relationship between intuition and metrics. Users are still required to interpret both independently.

4.1.4 Ranking-Based Interface

Another idea involved automatically ranking counterfactual explanations based on selected metrics.

Strengths:

- Reduces decision effort for users
- Provides a clear outcome based on defined criteria
- Supports quick identification of “best” explanations

Limitations:

- Removes the need for user interpretation
- Obscures trade-offs between different metrics
- Does not expose disagreement between human judgment and model evaluation

Relation to the core problem:

This concept bypasses the central issue by replacing user judgment with automated ranking. As a result, it does not support the investigation of discrepancies between intuition and metrics.

4.1.5 Guided Task Mode

The final concept introduced a structured interaction flow in which users are guided through a sequence of evaluation steps.

In this concept, users:

1. evaluate explanations based on visual inspection
2. estimate evaluation metrics
3. are presented with actual model metrics
4. reflect on differences between their judgment and the model output

Strengths:

- Structures the evaluation process
- Separates intuition from metric-based evaluation
- Enables direct comparison between user input and model output
- Reduces cognitive load through step-by-step interaction

Limitations:

- Less flexible than free exploration dashboards
- Requires careful design to avoid over-guidance

Relation to the core problem:

This concept directly addresses the central challenge by making the relationship between intuition and metrics explicit. It operationalizes the gap identified in the empathize phase and creates a measurable interaction.

4.2 4.2 COCD Box — Structuring Divergence and Convergence

To support convergence, the generated ideas were mapped onto a COCD (Creative Problem Solving) box, positioning them based on feasibility and originality.




Figure 8: COCD box showing divergence and convergence of design ideas

The COCD box distinguishes four quadrants:

- Now: high feasibility, low originality
- How: high feasibility, higher originality
- Wow: high originality, lower feasibility
- Non-ideas: low feasibility and low relevance

Most ideas were positioned in the “Now” quadrant, including *Metrics First Dashboard*, *Case First Comparison Dashboard*, *Slice & Dice Method Comparison*, *The Lineup*, and *The Gap Game*. These concepts are feasible and align with established design patterns, but do not fundamentally change how users interpret explanations.

Several concepts were positioned in the “How” quadrant, including *The Footprint*, *The Mirror*, and the *Guided Task Mode Dashboard*. These ideas introduce new interaction patterns while remaining feasible within the scope of the project.

Two concepts, *The Tension Map* and *The Guided Task Game*, were placed in the “Wow” quadrant. While

these ideas offer high originality, they were considered less feasible to implement within the available time constraints.

The distribution of ideas indicates that initial ideation tended to gravitate toward familiar solutions. The COCD mapping made it possible to explicitly identify and select concepts that introduce meaningful innovation while remaining implementable.

4.3 4.3 Convergence — Selection of Design Direction

Based on the COCD analysis and the insights from the empathize and define phases, the *Guided Task Mode Dashboard* was selected as the primary design direction.

This decision is based on three considerations.

4.3.1 Alignment with User Needs

The guided task mode addresses the observed difficulties in interpreting counterfactual explanations by structuring the evaluation process. It reduces cognitive complexity and supports users in forming judgments step by step.

Unlike free exploration approaches, it provides guidance without removing user agency.

4.3.2 Mitigating Metric Bias

The empathize phase highlighted the risk of metric-driven bias. The guided task mode addresses this by delaying the presentation of evaluation metrics.

Users first form an independent judgment, after which they are exposed to model-based evaluation. This supports a clearer distinction between intuition and metrics.

4.3.3 Making the Core Problem Observable

The guided task mode makes the central problem measurable. By capturing both user judgments and model outputs, it enables the analysis of discrepancies between the two.

Alternative concepts either prioritize exploration or automate evaluation, but do not make this relationship explicit.

4.4 4.4 Design Direction

The selected design direction can be summarized as follows:

- a step-by-step interaction flow that guides users through evaluation tasks
- separation between intuitive judgment and metric-based evaluation
- explicit comparison between user input and model output
- structured reflection on discrepancies

This direction ensures that the design is directly aligned with the research question and provides a clear foundation for the prototype phase, where the guided interaction is translated into a concrete interface.

5 Prototype Phase

5.1 Prototype Description

The prototype is an interactive web-based dashboard built with Streamlit, designed to serve two types of users: **study participants** (data science students who evaluate counterfactual explanations) and **researchers** (who collect and analyse the resulting data). The dashboard runs on real data, PyTorch image tensors generated by five CF methods across 1 dataset (MNIST) and supports genuine interaction at every step, making it appropriate for a real user study rather than a demonstration.

The current application in this repository is organized into four user-facing pages in the sidebar:

- Guided Task Mode
- Mini Game
- Leaderboard
- Guided Task Results

Guided Task Mode remains the central research flow. Participants complete a locked seven-step protocol for one sampled case, and each step captures a different judgment type such as visual assessment, scalar estimates, forced choices, and written reflection. The lock is intentional so that human intuition is recorded before metrics are revealed.

The Mini Game and Leaderboard support engagement and repeated interaction with CF validity judgments. The Guided Task Results page aggregates completed guided sessions for analysis and export.

5.1.1 Code Structure in `src/`

The prototype was refactored into modular Python files under `src/` instead of one monolithic script.

- `src/app.py` handles page routing, global session state, and sidebar navigation.
- `src/intro.py` contains onboarding and participant entry for the guided protocol.
- `src/task.py` contains the full seven-step guided-task logic and persistence hooks.
- `src/minigame.py` implements the mini game loop and scoring.
- `src/leaderboard.py` renders ranking views for mini game sessions.
- `src/results.py` renders aggregate guided-task analytics and CSV export.
- `src/components.py` centralizes reusable UI components, metric explainers, and plotting helpers.
- `src/data_utils.py` centralizes CSV loading, tensor loading, metric retrieval, and dataset utilities.

This separation is deliberate. It improves maintainability, makes each page independently testable, reduces merge conflicts, and keeps study logic clearly separated from shared rendering and data access concerns.

5.1.2 Result Storage in `Game Results/`

For reporting and analysis, session outputs are stored as JSON logs in `Game Results/`.

- `Game Results/guided_task_results.json` stores one complete record per finished guided-task session.
- `Game Results/minigame_log.json` stores one complete record per finished mini game run, including per-round history.

These files are used for descriptive analysis, comparison of human choices against objective metrics, and CSV export from the results interfaces.

5.2 Design Decisions

5.2.1 Linear, locked step progression

The seven steps follow a strict linear order and participants cannot navigate freely between them. This decision was driven directly by the core research question: to measure how human judgment of CF quality diverges from objective evaluation metrics. If participants could jump immediately to a step showing metrics, they would anchor their judgments on those numbers rather than forming an independent opinion. The linear lock ensures that intuition is always captured *before* information is revealed, making the collected data meaningful for comparison. Each step builds on the previous one in a deliberate sequence: first pure visual judgment, then estimation, then selection, then revelation, then reflection.

5.2.2 Showing one random CF method in steps 1 and 2 before the full lineup

In steps 1 and 2, participants see only the original image and one randomly chosen CF method, without seeing the method identity. All five methods are shown together only in step 3. This forces an initial independent judgment and reduces method-name anchoring effects.

5.2.3 Combining free-text and structured input

The protocol combines structured responses (radio buttons, sliders, fixed choices) with mandatory reflection fields at key points. Structured fields support aggregation and comparison across participants, while free text captures reasoning that explains *why* participants diverge from metric-based rankings.

5.2.4 In-session guidance and explainability support

The interface includes a persistent step-specific information panel so participants can request context at any point in the session. This panel explains what the current step asks from the user, how to interpret the task objective, and how to read the metric terminology introduced in the onboarding. This was added to reduce confusion during longer sessions and to keep participant interpretation consistent without forcing users to return to the introduction page.

5.2.5 Hiding metrics until step 4

Evaluation metrics are first shown in step 4, after visual judgment and method selection have already been recorded. This timing protects the internal validity of the study by reducing metric-driven bias in earlier human judgments.

When metrics are revealed, the interface explicitly compares participant estimates with model values. It reports both the participant estimate and the actual metric outcome, and it labels model-side values using qualitative bands such as good, moderate, and poor where applicable. For validity, the system also checks agreement between the participant direction of judgment and the model outcome.

The estimation interface intentionally uses a 0 to 1 human slider for both validity and plausibility in step 2. This is easy for participants to interpret and keeps input consistent across users. In contrast, the underlying model metrics are not on a shared human scale. Validity is binary and maps cleanly to 0 or 1, while plausibility-related metrics such as IM1 can exceed 1. The separation between human-scale input and model-scale output is therefore explained explicitly in the interface and in step 4 to avoid misinterpretation.

5.2.6 The Game as the primary data collection mechanism

Step 3 captures the key comparative choice: which CF method the participant judges as best before seeing objective scores. This supports direct comparison between perceived quality and metric-defined quality across sessions.

5.2.7 Collecting all inputs per session, not per step

Rather than storing partial rows after each interaction, the guided-task log stores complete sessions after step 7. This keeps each entry analytically coherent and allows within-session comparison across all steps for the same participant.

5.2.8 Dataset choice

Only MNIST is used in this study, as handwritten digit recognition provides a simpler and more interpretable visual domain for participants. Digits are easier to judge intuitively, as it is clearer whether a modified image still looks like a recognisable digit. This reduces noise in the human judgments and keeps the focus on the validity and plausibility gap rather than on image complexity. Cases are sampled randomly from MNIST instances that have non-blank CF images available across a sufficient number of methods.

For the full Prototype final version screenshots we reference to the appendix [Section 10](#)

5.2.9 Archive

Earlier prototype version

The dashboard was initially built as a simpler Guess Mode interface in which participants were shown one CF at a time and asked only whether they thought it was valid (yes or no), with their response immediately compared against the actual correctness score. This version collected minimal data and did not capture the reasoning behind judgments or allow comparison across methods within the same case.

[Screenshot: Earlier Guess Mode interface]

The Guided Task Mode was developed as a replacement to address the limitations of this approach: the binary guess format did not allow measurement of plausibility perception, did not capture why participants made their choices, and did not enable cross-method comparison within a single case, all of which are necessary to answer the research question with depth.

Dashboard file: `app.py` (run with `streamlit run app.py`)

Data log: `game_log.json` (auto-generated, appended to after each completed session)

6 Test Phase (4 pages)

6.1 Research Question

6.2 Method

6.3 Results

6.4 Insights

7 Conclusions & Recommendations (4 pages)

7.1 Conclusions from User Testing

7.2 Improvement Recommendations

7.3 Final Design Concept

7.4 User Interaction in Context

8 Design Archive Description (0.5 page)

8.1 Archive Overview

This archive documents the first complete front-end prototype we built for the XAI study: an interactive Streamlit dashboard that translated our evaluation pipeline into a participant-facing experience. The goal of this version was to make the study flow testable from end to end and to verify that participants could compare counterfactual explanations in a structured way.

The prototype was implemented in `Dashboard/app.py`, with data and image handling in `Dashboard/data_utils.py`. It combined model outputs from MNIST and CIFAR-10, loaded original and counterfactual images, and guided participants through a step-by-step interaction sequence: visual judgment, metric estimation, method comparison, and final reflection. This structure was useful because it let us observe how judgments changed before and after metric information was revealed.

A key contribution of this V1 frontend was the integration of quantitative and qualitative feedback in one flow. Participants gave confidence and plausibility estimates, selected the method they considered best, and provided short written explanations for their choices. After each session, responses were automatically stored in `game_log.json`, which enabled efficient analysis without manual data collection. The dashboard also included a results page for quick aggregation checks (e.g., method preferences and confidence distributions), which helped us validate that logging and comparison logic worked during early testing.

At the same time, this archived version had usability limitations that motivated the redesign. Although metrics were shown, users often did not understand what those metrics meant, why they were included, or how to interpret high and low values. In addition, the interaction flow was not always intuitive: participants were asked in later steps (such as Step 7) to reflect on early choices from Step 1, but without enough in-context reminders this felt less user-friendly.

The newer frontend integrated in this report addresses these issues directly. We added clearer explanations of each metric (what it represents, why it matters, and how to interpret it), improved page-level instructions, and introduced repeated context and information cues so users consistently understand what they need to do on each screen. In that sense, the V1 archive remains an important bridge: it established the core interaction logic and data pipeline, while also making clear which design improvements were necessary for a better user experience. For the full prototype archive see figures 1 to 8 in Section 11

9 References

10 Appendix Final Version Prototype




Figure 9: HomeScreen 1




Figure 10: Homescreen 2

STEP 1 OF 7

Visual Inspection

What am I judging here?

Look at the images below. No metric information is shown yet. Use only your eyes.

Original: 5

Counterfactual (target 7)

Dataset: MNIST - Instance: 24 - Target: 7

Based on visual inspection only, does the counterfactual look successful?

A successful CF should look like it clearly belongs to the target class.

Your judgment:

☐ Yes, successful
 ☐ No, not successful
 ☒ Uncertain

Why do you think so? (optional but valuable)

e.g. The digit looks like a 3 but the stroke is a bit off...

Next >

Figure 11: Visual Inspection FV

STEP 2 OF 7

Your Prediction

What do validity and plausibility mean here?

Validity (your slider): how confident are you that the model was fooled? 0 = definitely not, 1 = definitely yes.

Plausibility (your slider): how realistic does the CF look to you as a human? 0 = very distorted or unnatural, 1 = looks like a real digit.

These are your personal estimates on a simple 0 to 1 human scale. The actual model metrics use different scales and will be revealed in step 4.

Now estimate the evaluation metrics for this counterfactual, based on what you see.

Original: 5

Counterfactual (target 7)

Rate the following based on what you see (0 to 1):

These are your human judgments on a simple 0-1 scale. In step 4 you will see the actual model metrics, note that those use different scales and directions, which we will explain clearly at that point.

Validity: do you think the model was fooled by this CF?

0 = definitely not, 1 = definitely yes

Your validity estimate

0.50

Plausibility: how realistic does this CF look to you?

0 = very unrealistic / distorted, 1 = very realistic / natural

Your plausibility estimate

0.50

Next >

Figure 12: User Prediction FV

Select the CF method you think is best:

☒ PIECE
 ☐ Min-Edit
 ☐ C-Min-Edit
 ☐ alibi-Proto-CF
 ☐ alibi-CF

How confident are you in this choice? (%)

60

Tell us why you made this choice, this is required.

Why did you pick this method? Were you surprised by any of the images? What stood out to you compared to the other methods?

e.g. PIECE looked the most natural, the digit was still recognisable but clearly shifted towards the target. The alibi methods looked too noisy...

Please explain your choice above before continuing.

Submit choice >

Figure 13: The Game




Figure 14: Actual metrics 1





Figure 15: Actual metric 2



actual metric 3}



Figure 16: Explanation & Feedback FV 1




Figure 17: Explanation & Feedback 2

STEP 6 OF 7 Deploy

Compare All Methods on This Case

How do I use this comparison table?

Look at each column (method) and each row (metric). For each metric, lower is better except for validity where 1 = valid is best. Use the three dropdowns to select which method you think wins on each criterion. You can pick the same method for multiple criteria or different ones.

In step 5 you saw that the overall best method was determined by combining validity and plausibility. But does the best overall method also win on every individual metric? Look at the table below and find out.

Reminder: in step 3 you picked **C-Min-Edit** as the best method.

Below is the image grid from step 3, same images, now with the full metric table.

What do these metrics mean? (click to expand)

Figure 18: Compare methods FV 1

What do these metrics mean? (click to expand)

What do these metrics mean?

- Validity (Correctness)**
Did the model actually predict the target class for this CF? 1 = Yes (valid), 0 = No (invalid). This is a binary pass/fail check, it says nothing about how the image looks, only whether the model was fooled.
- Plausibility (IM1)**
Does the CF look like a realistic image? IM1 measures how far the CF is from the training data distribution. LOWER IM1 = MORE plausible (closer to real data). Higher IM1 = the image looks unnatural or distorted.
- Implausibility score**
A combined score that captures how implausible the CF looks overall. HIGHER = more implausible (looks worse). This is sometimes confused with plausibility, remember: high implausibility is bad.
- IM2 / L2 distance**
Measures how much the CF image differs from the original in terms of pixel values. A lower score means the CF required fewer changes to fool the model, which is better. A higher score means large parts of the image were altered.

Full evaluation metrics per method:

	Validity	IM1 (lower=better)	Implausibility (lower=better)	L2 distance (lower=better)
PIECE	X Invalid	1.000	0.470	23.205
Min-Edit	✓ Valid	1.045	0.482	22.081
C-Min-Edit	✓ Valid	0.931	0.483	21.848
alibi-Proto-CF	✓ Valid	0.587	0.141	22.905
alibi-CF	X Invalid	N/A	N/A	N/A

Figure 19: Compare methods 2 FV

Per-metric winners from the table above:

Most valid methods Min-Edit, C-Min-Edit, alibi-Pr...	Best IM1 (plausibility) alibi-Proto-CF (0.587)	Best implausibility alibi-Proto-CF (0.141)	Best L2 distance C-Min-Edit (21.848)
--	--	--	--

The overall best method from step 5 was **alibi-Proto-CF**. Notice whether it wins on every individual metric or only on the combined score. This is the key insight: 'best overall' does not always mean 'best at everything'.

One question, this is required.

Knowing the full metric breakdown, would you change your step 3 pick of **C-Min-Edit**?

☒ No, I would still pick the same method

☐ Yes, I would change my pick

Explain your reasoning, required. What does the table tell you that changed (or confirmed) your view?

e.g. I would still pick alibi-Proto-CF because it wins on implausibility and L2, even though PIECE has a better IM1 score but is invalid...

Please explain your reasoning above before continuing.

Next →

Figure 20: Compare methods FV 3




Figure 21: CF & Thoughts FV




Figure 22: CF Thoughts 2




Figure 23: Mini-Game Homescreen




Figure 24: Minigame Information



Figure 25: Round Example 1



Figure 26: Round Example 2




Figure 27: Mini-Game Choice




Figure 28: Leaderbord

11 Appendix Archive Prototype




Figure 29: Homescreen




Figure 30: Visual Inspection




Figure 31: User Prediction




Figure 32: The Game




Figure 33: Actual Metrics




Figure 34: Explanation & Feedback




Figure 35: Compare methods 1.1

Amershi, Saleema, Daniel Weld, Mihaela Vorvoreanu, et al. 2019. “Guidelines for Human-AI Interac-




Figure 36: Compare methods 1.2




Figure 37: Archive CF & FT

- tion.” In *Proceedings of the 2019 CHI Conference*, 1–13. <https://www.microsoft.com/en-us/research/publication/guidelines-for-human-ai-interaction/>.
- Doshi-Velez, Finale, and Been Kim. 2017. “Towards a Rigorous Science of Interpretable Machine Learning.” *arXiv Preprint arXiv:1702.08608*. <https://arxiv.org/abs/1702.08608>.
- Guidotti, Riccardo, Anna Monreale, Salvatore Ruggieri, Franco Turini, Fosca Giannotti, and Dino Pedreschi. 2018. “A Survey of Methods for Explaining Black Box Models.” *ACM Computing Surveys* 51 (5): 1–42. <https://arxiv.org/abs/1802.01933>.
- Kaur, Harmanpreet, Harsha Nori, Samuel Jenkins, and Rich Caruana. 2020. “Interpreting Interpretability: Understanding Data Scientists’ Use of Interpretability Tools for Machine Learning.” In *Proceedings of the 2020 CHI Conference*, 1–14. <https://arxiv.org/abs/2001.10996>.
- Wachter, Sandra, Brent Mittelstadt, and Chris Russell. 2017. “Counterfactual Explanations Without Opening the Black Box: Automated Decisions and the GDPR.” *Harvard Journal of Law & Technology* 31 (2): 841–87. <https://arxiv.org/abs/1711.00399>.